

Statistical learning

Shrikrishna Anand Narvekar

Assistant Professor

Agnel Institute of Technology and Design

Content

- Introduction
- How we learn rules
- Descriptive vs Inferential statistics

Introduction

- Statistical analysis saw the development of two branches in the 18th century-
 - Bayesian Statistics(Believer):Uses **Prior Belief + New Data** to reach a conclusion.
 - Classical Statistics(Observer):Uses **Only the Current Data**.
- Both deal with probability, but in different ways.

How we learn rules

- **Deductive (Top-Down):** Start with a **Rule**, apply it to **Data**.
- Logic: “All humans need water” → “John is human” → “John needs water”
- **Inductive (Bottom-up):** Start with a **Rule**, apply it to **Data**.
- *The "Baby" Example:* Touching a hot stove + hot soup + hot iron = "Don't touch hot things!"

Descriptive vs. Inferential Statistics

Type	Purpose	Example
Descriptive	Summarizes the Past .	A fitness app showing you walked 10,000 steps yesterday.
Inferential	Predicts the Future/Whole .	Asking 1,000 voters their opinion to predict how 30 million will vote.

Finding signals in the Mess

- **What is "Noise"?** Random errors or irrelevant data that distract from the pattern.
- **Signal:** Your friend's voice (the data you want).
- **Noise:** The loud music and crowd shouting (the distraction).
- **Goal:** Machine Learning acts like "Noise-Cancelling Headphones"—it filters out the music to hear the message.

Explaining noise in terms of data

House Size (sq.ft)	Actual Price (₹ Lakhs)
--------------------	------------------------

1000	50
------	----

1200	60
------	----

1500	75
------	----

1800	90
------	----

2000	500
------	-----

- First four rows show a clear pattern:

Bigger house → Higher price 

- But last row:

2000 sq.ft → ₹500 Lakhs

This value is unusually high and may be:

- Data entry error
- Luxury villa in different location
- Typing mistake
- This abnormal value is called **Noise (Outlier)**.

Hard vs. Soft Modelling

- **Hard Computing (Analytical):** Following a strict, 100% exact recipe
- **Rule:** If $A + B = C$ then C is always same
- **Example :** When you type 567×432 , the computer uses a "hard" mathematical rule. There is no guesswork; the answer is either right or wrong.
- **Soft modelling:** Learning patterns to make a "best guess."
- **Rule:** Based on thousands of examples, this is probably the answer.
- **Example: : Netflix Recommendations.** There is no "hard" rule that says if you liked *Stranger Things*, you will definitely like *Wednesday*.

Dealing with messy data

- Data is rarely perfect. We have to manage three main problems:
 1. **Noise:** Irrelevant "jitter" or extra info that hides the truth.
 2. **Missing Values:** Gaps where data wasn't recorded.
 3. **Inconsistencies:** Data that contradicts itself.

Employee ID	Age	Experience (Years)	Department	Salary (₹ Lakhs)
E101	25	2	Sales	4.5
E102	30	5	Marketing	6.8
E103	29	4	Sales	150.0
E104	—	3	HR	5.2
E105	32	12	Intern	3.0
E106	28	4	Sales	6.0

Sample Space & Sample Points

- **The Sample Space (D):** This is the "world" of all possible outcomes. In Machine Learning, your entire dataset (the **Data Matrix**) is your Sample Space.
- **Example:** If you are analysing a deck of cards, the Sample Space is all 52 cards.
- **Sample Points (x):** This is a single "event" or a single row in your data. It is one specific instance from the Sample Space.
- **Example:** Drawing the 7 of Hearts is a "sample point." In a medical study, one single patient's records would be a sample point

Random Variables

- We call the characteristics of our data "Random Variables" because we don't know the value until we actually look at a specific example.
- These are the **Columns** in your table (Attributes).
- **The "House Hunting" Example:** If you are looking for a house, the **Random Variables** are:
 - Number of Bedrooms
 - Square Footage
 - Price



- **The “Zero Probability” Problem**

- Since there are infinite possible values: The probability of getting **exactly one specific value** is ZERO.
- Instead of asking: What is probability of exactly 20 minutes?
- We ask:
- What is probability that commute time is between 18 and 22 minutes?
- For continuous variables, we always find probability for a **range**.
- Instead of looking at a single point, we look at the probability of a value falling within a **range** (between V_{1x} and V_{2x})

$$\begin{aligned} P(v_{1x} \leq x < v_{2x}) &= \int_{v_{1x}}^{v_{2x}} p(x) dx \\ p(x) &\geq 0, \text{ and } \int_{-\infty}^{\infty} p(x) dx = 1 \end{aligned} \tag{3.3}$$

What is PDF

- A **Probability Density Function (PDF)** is:
- A smooth curve
- That shows how values are distributed
- Probability is the **area under the curve**
It looks like a smooth wave (not spikes like PMF).



How do we calculate Probability

- To find probability between two values:
- We calculate the **area under the curve** between those two points.
- Mathematically, we use **integration**.
- Simple meaning:
- Probability = Area under curve between V_1 and V_2

Rules of PDF

- **Rule 1: Density must be zero or positive**
- $p(x) \geq 0$
- You cannot have negative density.
- **Rule 2: Total area under curve must equal 1**
- The entire area under the curve = 1 (which means 100%)
- If total area $\neq 1 \rightarrow$ it is not a valid PDF.

- **Very Small Interval Idea (Δx Concept)**
- If we take a very small range:
- x_0 to $x_0 + \Delta x$
- Then,
- Probability $\approx$ Density at that point $\times$ Width of range
- In simple words:
- Small Probability $\approx$ Height $\times$ Width
- Just like area of a rectangle.

$$P(x_0 \leq x < x_0 + \Delta x) \approx p(x = x_0) \Delta x$$

- **Data Scientist perspective:**
- From a data point of view:
- Density means:

- How frequently values occur in a very tiny interval.
- If:
- N = total observations
 ΔN = number of values in small interval
- Then:
- Density = (Relative frequency) $\div$ (Width of interval)
- As interval becomes extremely small $\rightarrow$ it becomes PDF.

$$p(x) = \lim_{N \rightarrow \infty, \Delta x \rightarrow 0} \left[\frac{\Delta N / N}{\Delta x} \right]$$

Descriptive Measures of Probability Distributions

- One of the most important descriptive values of a random variable is the point around which distribution is centered, known as measure of **central tendency**.
- The expected (mean or average) value is the best measure for central tendency.
- Expected value is just the **average** of a random situation.

- **Example 1: Dice Roll (Discrete Case)**

- Suppose you roll a fair dice.

- Possible outcomes:

1, 2, 3, 4, 5, 6

Each has probability = $1/6$

- Expected value:

- $E[X] = 1(1/6) + 2(1/6) + 3(1/6) + 4(1/6) + 5(1/6) + 6(1/6)$

- $= 3.5$

- You will **never actually get 3.5**

But if you roll many times, the average will approach 3.5.

- **For Discrete Data (Countable steps, like a dice roll):** You use simple addition. You multiply each outcome by its probability and sum them up.

- $$E[X] = \mu = \sum_{x \in S} xp(x)$$

- where $P(x)$ is probability mass function

- **Example 2: Rainfall (Continuous Case)**

- Suppose rainfall can be any value like:

- 12.2 mm, 12.25 mm, 12.251 mm...

- Here values are continuous.

- So instead of adding, we use **integration**.

- But concept is same:

- Multiply each value by how likely it is
Add everything together

- That gives the average rainfall.

- **For Continuous Data (Smooth curves):** You use calculus (integration). You multiply every possible value (x) by its likelihood (p(x)) and add them all up.

- $$E [x] = \mu = \int_{-\infty}^{\infty} xp(x)dx$$

- where p(x) is the probability density function.

Numerical

- A data scientist is evaluating a simple "Recommender System" for a streaming app. The random variable x represents the number of stars (1 to 4) a user is likely to give a specific genre based on historical data.

- | x | $P(x)$ |
|-----|--------|
| 1 | 0.1 |
| 2 | 0.3 |
| 3 | 0.4 |
| 4 | 0.2 |

Questions: Verify that this is valid PMF

Calculate Expected Value $E[x]$

(Average Rating)

Find Variance and Standard deviations

- A food delivery company tracks number of orders received per hour. The following data was collected:

Orders per hour (x)	Frequency
0	5
1	10
2	20
3	10
4	5

- **Questions:**
- Construct the probability mass function $P(x)$.
- Verify that it is a valid PMF.
- Find the expected number of orders per hour.
- Find variance and standard deviation.

Important properties

- **1. Constants stay the same:** The "average" of a constant number (like 5) is always 5.
 $E[C] = C$; $C = \text{constant}$
- **Example:** If you always get ₹500 salary bonus: $E[C] = C$
- Average of 500 is 500.
- **2. Scaling:** If you multiply all your data by 10, your average also multiplies by 10.
- $E[Cx] = C E[x]$
- **Example:** If daily salary is ₹1000
Now company doubles it $\rightarrow$ ₹2000
- New average also doubles.

- **3. Additivity:** The average of a bunch of things added together is the same as adding their individual averages.
- $E[\sum_{j=1}^n x_j] = \sum_{j=1}^n E[x_j]$
- **4. Joint Variables:** When dealing with two variables at once (like height and weight), you use their "joint" probability to find the average of their product.
- If there are two random variables x_1 and x_2 defined on a sample space, then
- $E[x_1 x_2] = \iint_{-\infty}^{\infty} x_1 x_2 p(x_1, x_2) dx_1 dx_2$
- **Why Do We Need Joint Probability?**
- Because now two things are happening together.
- If: X_1 =Height, X_2 =Weight
- Height and weight are related.
So we must look at their **joint behaviour** — how they occur together

5. Functions: If you want the average of x^2 you just plug x^2 into the formula where x used to be.

- $E[f(x)] = \mu = \int_{-\infty}^{\infty} f(x)p(x)dx$

- **Vectors:** If your data is a list of numbers (a vector) rather than a single value, the math works exactly the same way—you just do it for the whole list at once.

- For a vector random variable $\mathbf{x}$,

- $$E[\mathbf{x}] = \mu = \int_{-\infty}^{\infty} \mathbf{x}p(\mathbf{x})d\mathbf{x}$$

- In discrete case,

- $$E[\mathbf{x}] = \mu = \sum_{\mathbf{x} \in S} \mathbf{x}p(\mathbf{x})$$

-

Numerical

- A multi-specialty hospital has deployed a **Machine Learning-based patient triage system**. The ML model predicts patient severity and dynamically allocates doctors. After analysing historical data, the **data science team** observes that the waiting time (in minutes) for patients follows a probability density function:
 - $p(x) = kx$ for $0 \leq x \leq 10$
 1. Find Value of k
 2. Find probability patient waits between 2 and 5 minutes
 3. Find Expected waiting time

Numerical

- An autonomous vehicle uses a machine learning–based distance estimation model to predict the distance X (in meters) to nearby objects. The predicted distance follows a continuous probability density function:
- $p(x) = 0.1$ for $0 \leq x \leq 10$
- Assuming the model output follows this distribution, compute the probability that an object is detected between **2 and 5 meters** from the vehicle.

Variance

- The expected value (mean) describes the central tendency of a random variable.
- We also need to measure the spread of values around the mean.
- Variance is desired to know the ‘spread’ of the random variable about the mean.
- It is denoted as $\text{Var}[X]$ or σ^2

Variance (Continuous Case)

- $\text{Var}[x] = \sigma^2 = E[(x - \mu)]^2 = \int_{-\infty}^{\infty} (x - \mu)^2 p(x) dx$
- It is the weighted average of squared deviations
- from the mean.

Variance (Discrete Case)

- $\text{Var}[X] = \sigma^2 = E[(X - \mu)^2]$
- Summation Form:
- $\text{Var}[X] = \sum (x - \mu)^2 P(x), \quad x \in S$
- It represents weighted squared deviations
- from the mean.

Derivation of Variance Formula

- $\text{Var}[X] = E[(X - E[X])^2]$
- $= E[X^2 - 2X E[X] + (E[X])^2]$
- Using linearity of expectation:
- $= E[X^2] - (E[X])^2$

-
- Final Formula:
 - $\text{Var}[X] = E[X^2] - (E[X])^2$
 - This form is often easier to compute
 - in practical problems.